

SARSA Implementation- Treasure Hunter Game

Code for Testing:

```
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from collections import defaultdict
import time
```

```
class TreasureHunterSARSA:
```

```
    """
```

```
    🏴‍☠️ A Treasure Hunter that learns using SARSA algorithm!
```

The agent starts at the bottom-left and must find the treasure while avoiding dangerous traps in a grid world.

```
    """
```

```
    def __init__(self, grid_size=5, alpha=0.1, gamma=0.9, epsilon=0.1):
```

```
        """
```

```
        Initialize our Treasure Hunter
```

```
        Args:
```

```
            grid_size: Size of the grid world (grid_size x grid_size)
```

```
            alpha: Learning rate (how fast we learn from mistakes)
```

```
            gamma: Discount factor (how much we care about future rewards)
```

```
            epsilon: Exploration rate (how often we try random actions)
```

```
        """
```

```
        self.size = grid_size
```

```
        self.alpha = alpha
```

```
        self.gamma = gamma
```

```
        self.epsilon = epsilon
```

```
        # Actions: 0=UP, 1=RIGHT, 2=DOWN, 3=LEFT
```

```
        self.actions = ['UP', 'RIGHT', 'DOWN', 'LEFT']
```

```
        self.action_effects = [(-1, 0), (0, 1), (1, 0), (0, -1)]
```

```
        # Initialize Q-table: Q[state][action] = expected reward
```

```
        self.Q = defaultdict(lambda: np.zeros(4))
```

```
        # Game setup
```

```
        self.start_pos = (grid_size-1, 0) # Bottom-left corner
```

```
        self.treasure_pos = (0, grid_size-1) # Top-right corner
```

```
        self.traps = [(1, 1), (2, 3), (3, 2)] # Dangerous traps!
```

```

# Keep track of learning progress
self.episode_rewards = []
self.episode_steps = []

def get_state_key(self, pos):
    """Convert position tuple to string key for Q-table"""
    return f'{pos[0]},{pos[1]}'

def is_valid_position(self, pos):
    """Check if position is within grid boundaries"""
    row, col = pos
    return 0 <= row < self.size and 0 <= col < self.size

def get_reward(self, pos):
    """
    🎁 Reward function - what happens at each position?
    """
    if pos == self.treasure_pos:
        return 100 # 💰 Found the treasure!
    elif pos in self.traps:
        return -50 # 💣 Hit a trap!
    else:
        return -1 # ⌚ Small penalty for each step (encourages efficiency)

def take_action(self, pos, action):
    """
    🚶 Move the agent and return new position and reward
    """
    row, col = pos
    d_row, d_col = self.action_effects[action]
    new_pos = (row + d_row, col + d_col)

    # If move goes outside grid, stay in current position
    if not self.is_valid_position(new_pos):
        new_pos = pos

    reward = self.get_reward(new_pos)
    return new_pos, reward

def choose_action_epsilon_greedy(self, state_key, epsilon=None):
    """
    🎲 ε-greedy policy: Sometimes explore, sometimes exploit

```

```

"""
if epsilon is None:
    epsilon = self.epsilon

if np.random.random() < epsilon:
    # 🎯 Explore: Choose random action
    return np.random.choice(4)
else:
    # 🧠 Exploit: Choose best known action
    return np.argmax(self.Q[state_key])


def choose_action_boltzmann(self, state_key, temperature=1.0):
    """
    🎲 Boltzmann (Softmax) policy: Probabilistic action selection
    """
    q_values = self.Q[state_key]
    # Avoid overflow by subtracting max
    exp_values = np.exp((q_values - np.max(q_values)) / temperature)
    probabilities = exp_values / np.sum(exp_values)
    return np.random.choice(4, p=probabilities)

def choose_action_greedy(self, state_key):
    """
    🎯 Pure Greedy policy: Always choose best action (no exploration)
    """
    return np.argmax(self.Q[state_key])

def choose_action(self, state_key, policy='epsilon_greedy'):
    """
    🚗 Action selection based on chosen policy
    """
    if policy == 'epsilon_greedy':
        return self.choose_action_epsilon_greedy(state_key)
    elif policy == 'boltzmann':
        return self.choose_action_boltzmann(state_key)
    elif policy == 'greedy':
        return self.choose_action_greedy(state_key)
    else:
        raise ValueError(f"Unknown policy: {policy}")

def run_episode(self, policy='epsilon_greedy', max_steps=100):
    """

```

 Run one complete episode using SARSA algorithm


Returns:

total_reward: Total reward collected in this episode

steps: Number of steps taken

path: List of positions visited

"""

 Start at the beginning

current_pos = self.start_pos


current_state_key = self.get_state_key(current_pos)

current_action = self.choose_action(current_state_key, policy)


total_reward = 0

steps = 0

path = [current_pos]


 Keep moving until we reach treasure, trap, or max steps

while steps < max_steps:

 Take the action

next_pos, reward = self.take_action(current_pos, current_action)

next_state_key = self.get_state_key(next_pos)


 Choose next action using the same policy

next_action = self.choose_action(next_state_key, policy)


 SARSA Update: Learn from this experience!

current_q = self.Q[current_state_key][current_action]

next_q = self.Q[next_state_key][next_action]

 The SARSA magic happens here!


```
self.Q[current_state_key][current_action] += self.alpha * (
    reward + self.gamma * next_q - current_q
)
```

 Update tracking variables

total_reward += reward

steps += 1

path.append(next_pos)

 Check if episode is over

if next_pos == self.treasure_pos or next_pos in self.traps:

break

```

# 🔄 Move to next state and action
current_pos = next_pos
current_state_key = next_state_key
current_action = next_action

return total_reward, steps, path

def train(self, episodes=1000, policy='epsilon_greedy', verbose=True):
    """
    🎓 Train the agent for multiple episodes
    """
    print(f"🚀 Training Treasure Hunter with {policy} policy...")
    print(f"📊 Running {episodes} episodes...\n")

    self.episode_rewards = []
    self.episode_steps = []

    for episode in range(episodes):
        reward, steps, path = self.run_episode(policy)
        self.episode_rewards.append(reward)
        self.episode_steps.append(steps)

        # 📉 Decay exploration over time
        if policy == 'epsilon_greedy' and episode > 0 and episode % 100 == 0:
            self.epsilon *= 0.95

        # 📈 Progress updates
        if verbose and (episode + 1) % 200 == 0:
            avg_reward = np.mean(self.episode_rewards[-100:])
            avg_steps = np.mean(self.episode_steps[-100:])
            print(f"Episode {episode + 1}: Avg Reward = {avg_reward:.2f}, "
                  f"Avg Steps = {avg_steps:.2f},  $\epsilon$  = {self.epsilon:.3f}")

    print("✅ Training completed!")

def test_policy(self, policy='greedy', num_tests=5):
    """
    🧪 Test the learned policy
    """
    print(f"\n🧪 Testing learned policy ({policy})...")

```

```

original_epsilon = self.epsilon
test_results = []

for test in range(num_tests):
    if policy == 'greedy':
        self.epsilon = 0 # No exploration during testing

    reward, steps, path = self.run_episode(policy)
    test_results.append((reward, steps, path))

    print(f"Test {test + 1}: Reward = {reward}, Steps = {steps}")

self.epsilon = original_epsilon # Restore original epsilon
return test_results

def visualize_grid(self, path=None, title="Treasure Hunter Grid"):
    """
    🎮 Visualize the grid world with optional path
    """
    fig, ax = plt.subplots(1, 1, figsize=(8, 8))

    # Create grid
    grid = np.zeros((self.size, self.size))

    # Mark special positions
    treasure_row, treasure_col = self.treasure_pos
    grid[treasure_row, treasure_col] = 3 # Treasure

    start_row, start_col = self.start_pos
    grid[start_row, start_col] = 1 # Start

    for trap in self.traps:
        trap_row, trap_col = trap
        grid[trap_row, trap_col] = 2 # Traps

    # Create color map
    colors = ['white', 'lightgreen', 'red', 'gold']
    from matplotlib.colors import ListedColormap
    cmap = ListedColormap(colors)

    # Plot grid
    im = ax.imshow(grid, cmap=cmap, vmin=0, vmax=3)

```

```

# Add path if provided
if path:
    path_rows = [pos[0] for pos in path]
    path_cols = [pos[1] for pos in path]
    ax.plot(path_cols, path_rows, 'b-', linewidth=3, alpha=0.7, label='Path')
    ax.plot(path_cols[0], path_rows[0], 'go', markersize=15, label='Start')
    ax.plot(path_cols[-1], path_rows[-1], 'ro', markersize=15, label='End')


# Add grid lines
ax.set_xticks(np.arange(-0.5, self.size, 1), minor=True)
ax.set_yticks(np.arange(-0.5, self.size, 1), minor=True)
ax.grid(which="minor", color="black", linestyle='-', linewidth=1)

# Labels and title
ax.set_title(title, fontsize=16, fontweight='bold')
ax.set_xlabel('Column', fontsize=12)
ax.set_ylabel('Row', fontsize=12)

# Legend
legend_elements = [
    plt.Rectangle((0,0),1,1, facecolor='lightgreen', label='Start'),
    plt.Rectangle((0,0),1,1, facecolor='gold', label='Treasure'),
    plt.Rectangle((0,0),1,1, facecolor='red', label='Trap')
]
if path:
    legend_elements.append(plt.Line2D([0], [0], color='blue', linewidth=3, label='Path'))

ax.legend(handles=legend_elements, loc='center left', bbox_to_anchor=(1, 0.5))

plt.tight_layout()
plt.show()

def visualize_learning_progress(self):
    """
     Show how the agent improved over time
    """
    fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(15, 5))

    # Smooth the curves for better visualization
    window = 50
    if len(self.episode_rewards) > window:
        smooth_rewards = np.convolve(self.episode_rewards, np.ones(window)/window,
mode='valid')

```

```

smooth_steps = np.convolve(self.episode_steps, np.ones(window)/window,
mode='valid')
episodes = np.arange(window-1, len(self.episode_rewards))
else:
    smooth_rewards = self.episode_rewards
    smooth_steps = self.episode_steps
    episodes = np.arange(len(self.episode_rewards))

# Plot rewards
ax1.plot(episodes, smooth_rewards, 'b-', linewidth=2)
ax1.set_title('🏆 Learning Progress: Rewards', fontsize=14, fontweight='bold')
ax1.set_xlabel('Episode')
ax1.set_ylabel('Average Reward')
ax1.grid(True, alpha=0.3)

# Plot steps
ax2.plot(episodes, smooth_steps, 'r-', linewidth=2)
ax2.set_title('👣 Learning Progress: Steps', fontsize=14, fontweight='bold')
ax2.set_xlabel('Episode')
ax2.set_ylabel('Average Steps')
ax2.grid(True, alpha=0.3)

plt.tight_layout()
plt.show()

def visualize_q_values(self):
    """
    🧠 Visualize the learned Q-values as a heatmap
    """
    fig, axes = plt.subplots(2, 2, figsize=(12, 10))
    action_names = ['UP ⬆️', 'RIGHT ➡️', 'DOWN ⬇️', 'LEFT ⬅️']

    for action_idx, ax in enumerate(axes.flat):
        # Create Q-value grid for this action
        q_grid = np.zeros((self.size, self.size))

        for i in range(self.size):
            for j in range(self.size):
                state_key = self.get_state_key((i, j))
                q_grid[i, j] = self.Q[state_key][action_idx]

        # Plot heatmap
        sns.heatmap(q_grid, annot=True, fmt='.1f', cmap='RdYlBu_r',

```



```

        center=0, ax=ax, cbar_kws={'label': 'Q-value'})
    ax.set_title(f'Q-values for {action_names[action_idx]}', fontweight='bold')

plt.suptitle('🧠 Learned Q-values for All Actions', fontsize=16, fontweight='bold')
plt.tight_layout()
plt.show()

def compare_policies(self, policies=['epsilon_greedy', 'boltzmann', 'greedy'],
                    episodes=500, num_tests=10):
    """
    🗡️ Compare different policies side by side
    """
    print("🗡️ POLICY COMPARISON BATTLE! 🗡️\n")

    results = {}

    for policy in policies:
        print(f"🤖 Training with {policy} policy...")

        # Reset Q-table for fair comparison
        self.Q = defaultdict(lambda: np.zeros(4))
        self.epsilon = 0.1 # Reset epsilon

        # Train the agent
        self.train(episodes, policy, verbose=False)

        # Test the trained policy
        test_results = self.test_policy('greedy', num_tests)

        # Calculate statistics
        test_rewards = [result[0] for result in test_results]
        test_steps = [result[1] for result in test_results]

        results[policy] = {
            'avg_reward': np.mean(test_rewards),
            'std_reward': np.std(test_rewards),
            'avg_steps': np.mean(test_steps),
            'std_steps': np.std(test_steps),
            'success_rate': sum(1 for r in test_rewards if r > 0) / len(test_rewards) * 100,
            'training_rewards': self.episode_rewards.copy()
        }

    print(f"✅ {policy}: Avg Reward = {results[policy]['avg_reward']:.1f} ")


```

```

f'(±{results[policy]['std_reward']:.1f}), "
f"Success Rate = {results[policy]['success_rate']:.1f}%\n")

# Visualize comparison
self.plot_policy_comparison(results, policies)

return results

def plot_policy_comparison(self, results, policies):
    """
     Plot comparison between different policies
    """
    fig, ((ax1, ax2), (ax3, ax4)) = plt.subplots(2, 2, figsize=(15, 10))

    # Colors for different policies
    colors = ['blue', 'red', 'green', 'orange', 'purple']

    # 1. Training curves
    for i, policy in enumerate(policies):
        training_rewards = results[policy]['training_rewards']
        window = 50
        if len(training_rewards) > window:
            smooth_rewards = np.convolve(training_rewards, np.ones(window)/window,
mode='valid')
            episodes = np.arange(window-1, len(training_rewards))
        else:
            smooth_rewards = training_rewards
            episodes = np.arange(len(training_rewards))

        ax1.plot(episodes, smooth_rewards, color=colors[i], linewidth=2,
            label=policy.replace('_', ' ').title())

    ax1.set_title('🏃 Training Progress Comparison', fontweight='bold')
    ax1.set_xlabel('Episode')
    ax1.set_ylabel('Average Reward')
    ax1.legend()
    ax1.grid(True, alpha=0.3)

    # 2. Average rewards comparison
    avg_rewards = [results[policy]['avg_reward'] for policy in policies]
    std_rewards = [results[policy]['std_reward'] for policy in policies]

    bars1 = ax2.bar(range(len(policies)), avg_rewards, yerr=std_rewards,

```

```

        color=colors[:len(policies)], alpha=0.7, capsize=5)
ax2.set_title(' 🏆 Average Test Rewards', fontweight='bold')
ax2.set_xlabel('Policy')
ax2.set_ylabel('Average Reward')
ax2.set_xticks(range(len(policies)))
ax2.set_xticklabels([p.replace('_', ' ').title() for p in policies])
ax2.grid(True, alpha=0.3)

# Add value labels on bars
for bar, value in zip(bars1, avg_rewards):
    ax2.text(bar.get_x() + bar.get_width()/2, bar.get_height() + 1,
             f'{value:.1f}', ha='center', va='bottom', fontweight='bold')

```

3. Average steps comparison

```

avg_steps = [results[policy]['avg_steps'] for policy in policies]
std_steps = [results[policy]['std_steps'] for policy in policies]

```

```

bars2 = ax3.bar(range(len(policies)), avg_steps, yerr=std_steps,
               color=colors[:len(policies)], alpha=0.7, capsize=5)
ax3.set_title(' 🏁 Average Steps to Goal', fontweight='bold')
ax3.set_xlabel('Policy')
ax3.set_ylabel('Average Steps')
ax3.set_xticks(range(len(policies)))
ax3.set_xticklabels([p.replace('_', ' ').title() for p in policies])
ax3.grid(True, alpha=0.3)

```

Add value labels on bars

```

for bar, value in zip(bars2, avg_steps):
    ax3.text(bar.get_x() + bar.get_width()/2, bar.get_height() + 0.5,
             f'{value:.1f}', ha='center', va='bottom', fontweight='bold')

```

4. Success rate comparison

```

success_rates = [results[policy]['success_rate'] for policy in policies]

```

```

bars3 = ax4.bar(range(len(policies)), success_rates,
               color=colors[:len(policies)], alpha=0.7)
ax4.set_title(' 🎯 Success Rate', fontweight='bold')
ax4.set_xlabel('Policy')
ax4.set_ylabel('Success Rate (%)')
ax4.set_xticks(range(len(policies)))
ax4.set_xticklabels([p.replace('_', ' ').title() for p in policies])
ax4.set_ylim(0, 100)
ax4.grid(True, alpha=0.3)

```

```
# Add value labels on bars
for bar, value in zip(bars3, success_rates):
    ax4.text(bar.get_x() + bar.get_width()/2, bar.get_height() + 1,
             f'{value:.1f}%', ha='center', va='bottom', fontweight='bold')

plt.suptitle('🗝️ SARSA Policy Comparison Results', fontsize=16, fontweight='bold')
plt.tight_layout()
plt.show()
```

🎮 INTERACTIVE DEMO FUNCTIONS

```
def demo_basic_sarsa():
    """ 🚀 Basic SARSA Demo """
    print("=" * 60)
    print(" 🏠 WELCOME TO TREASURE HUNTER SARSA! 🏠 ")
    print("=" * 60)
    print("Our brave hunter starts at bottom-left ( ● ) and must find")
    print("the treasure at top-right ( ● ) while avoiding traps ( ■ )!")
    print()

    # Create and train agent
    agent = TreasureHunterSARSA(grid_size=5, alpha=0.1, gamma=0.9, epsilon=0.3)

    # Show initial grid
    print(" 🗺️ Here's our treasure map:")
    agent.visualize_grid(title=" 🗺️ Treasure Hunter Map")

    # Train the agent
    agent.train(episodes=1000, policy='epsilon_greedy')

    # Show learning progress
    print("\n 📊 Let's see how our hunter learned:")
    agent.visualize_learning_progress()

    # Test the learned policy
    print("\n 🧪 Testing our trained treasure hunter:")
    test_results = agent.test_policy('greedy', num_tests=3)

    # Show the best path
    best_result = min(test_results, key=lambda x: x[1]) # Minimum steps
    print(f"\n 🏆 Best path found in {best_result[1]} steps with reward {best_result[0]}!")
```

```

agent.visualize_grid(path=best_result[2], title="🏆 Best Path Found")

# Show Q-values
print("\n🧠 Here's what our hunter learned (Q-values):")
agent.visualize_q_values()

def demo_policy_comparison():
    """🔪 Policy Comparison Demo"""
    print("=" * 60)
    print("🔪 POLICY BATTLE ARENA! 🔪")
    print("=" * 60)
    print("Let's see which policy performs best!")
    print()

    agent = TreasureHunterSARSA(grid_size=5)

    # Compare different policies
    results = agent.compare_policies(
        policies=['epsilon_greedy', 'boltzmann', 'greedy'],
        episodes=500,
        num_tests=10
    )

    # Announce the winner!
    winner = max(results.items(), key=lambda x: x[1]['success_rate'])
    print(f"🏆 THE WINNER IS: {winner[0].replace('_', ' ').title().upper()}!")
    print(f"Success Rate: {winner[1]['success_rate']:.1f}%")
    print(f"Average Reward: {winner[1]['avg_reward']:.1f}")

def interactive_parameter_tuning():
    """🔧 Interactive Parameter Tuning"""
    print("=" * 60)
    print("🔧 PARAMETER TUNING LAB! 🔧")
    print("=" * 60)
    print("Let's see how different parameters affect learning!")
    print()

    # Test different learning rates
    print("📊 Testing different learning rates...")
    alphas = [0.01, 0.1, 0.5]

    fig, axes = plt.subplots(1, 3, figsize=(18, 5))

```

```

for i, alpha in enumerate(alphas):
    agent = TreasureHunterSARSA(alpha=alpha, epsilon=0.2)
    agent.train(epochs=800, verbose=False)

    # Plot learning curve
    window = 50
    if len(agent.episode_rewards) > window:
        smooth_rewards = np.convolve(agent.episode_rewards, np.ones(window)/window,
mode='valid')
        episodes = np.arange(window-1, len(agent.episode_rewards))
    else:
        smooth_rewards = agent.episode_rewards
        episodes = np.arange(len(agent.episode_rewards))

    axes[i].plot(episodes, smooth_rewards, 'b-', linewidth=2)
    axes[i].set_title(f'Learning Rate  $\alpha$  = {alpha}', fontweight='bold')
    axes[i].set_xlabel('Episode')
    axes[i].set_ylabel('Average Reward')
    axes[i].grid(True, alpha=0.3)

plt.suptitle('🔧 Effect of Learning Rate on Training', fontsize=16, fontweight='bold')
plt.tight_layout()
plt.show()

print("📄 Observations:")
print(" • Low  $\alpha$  (0.01): Slow but stable learning")
print(" • Medium  $\alpha$  (0.1): Good balance of speed and stability")
print(" • High  $\alpha$  (0.5): Fast but potentially unstable learning")

# 🎯 MAIN EXECUTION
if __name__ == "__main__":
    print("🎮 Choose your adventure:")
    print("1. 🚀 Basic SARSA Demo")
    print("2. ✂️ Policy Comparison")
    print("3. 🔧 Parameter Tuning")
    print("4. 🎪 Full Demo (All of the above!)")

    choice = input("\nEnter your choice (1-4): ").strip()

    if choice == '1':
        demo_basic_sarsa()

```

```

elif choice == '2':
    demo_policy_comparison()
elif choice == '3':
    interactive_parameter_tuning()
elif choice == '4':
    print("🏠 RUNNING FULL DEMO!\n")
    demo_basic_sarsa()
    print("\n" + "="*60 + "\n")
    demo_policy_comparison()
    print("\n" + "="*60 + "\n")
    interactive_parameter_tuning()
else:
    print("🎮 Running Basic Demo...")
    demo_basic_sarsa()

print("\n🎉 explore SARSA! 🎉")

```

student:

Explore different SARSA policies and implement in the code

Environment Modifications:

Try different grid sizes agent = TreasureHunterSARSA(grid_size=8)

Larger world

Change trap positions agent.traps = [(2, 2), (4, 4)]

Different obstacles

Multiple treasures agent.treasure_pos = [(0, 4), (4, 0)]

Two goals! 2. ⚙️ Parameter Sensitivity Analysis:

Test learning rates for alpha in [0.01, 0.1, 0.5, 0.9]: agent =

TreasureHunterSARSA(alpha=alpha) agent.train(epochs=500)

Compare results # Test discount factors for gamma in [0.5, 0.8, 0.9, 0.99]: agent =

TreasureHunterSARSA(gamma=gamma) agent.train(epochs=500) 3. 🎯 Policy

Comparison:

Compare with Q-learning (add Q-learning method) # Test different epsilon decay schedules

Try UCB exploration policy and Interactive Features